

Kouventa Financial Systems

Employee Onboarding & Technical Guidelines

Department: Artificial Intelligence & Machine Learning Operations

Document Version: 1.4 (Updated August 2026)

Welcome to the AI/ML Operations team at Kouventa! Your role as an AI/ML Engineer is critical to our mission of building intelligent, secure, and highly efficient financial systems. This document outlines your core responsibilities, infrastructure access, and our strict security protocols.

1. Core Responsibilities & Workflow

As an engineer on this team, your primary workflow revolves around researching, developing, and evaluating Proof of Concepts (POCs) before they are merged into our production banking applications.

- POC Development:** All new models must be developed within the Kouv-Sandbox environment. You are expected to document your model's architecture, training data sources, and inference speed.
- Evaluation Metrics:** For classification tasks, we strictly require a minimum F1-score of 0.92 before a model can move to the staging environment. For embedding and RAG systems, context retrieval accuracy must be logged using the internal Eval-RAG script.
- Weekly Reporting:** Engineers must submit a brief technical summary of their POC progress every Friday by 3:00 PM WIB to the project lead.

2. IT Infrastructure & Hardware Protocols

Your development environment has been pre-configured to ensure compliance with banking security standards.

Resource	Address / Endpoint	Access Requirements
Model Registry	10.50.2.100/registry	VPN + YubiKey Authentication
Vector Database (Staging)	10.50.3.45:5432	Read-only access by default

Resource	Address / Endpoint	Access Requirements
GPU Cluster Allocation	gpu-alloc.kouventa.internal	Requires approved Jira Ticket (Max 24h lease)
Knowledge Management System	kms.kouventa.internal	Standard SSO

Note: Unauthorized installation of third-party Python packages outside of the approved requirements-strict.txt manifest is disabled on company hardware.

3. Data Privacy & POJK Compliance

Because we handle sensitive financial data, strict adherence to national banking regulations, specifically Otoritas Jasa Keuangan (POJK) guidelines, is mandatory.

- **Policy #KVT-SEC-902 (Data Anonymization):** Never use raw, unmasked customer data for local model training or POC testing. All datasets must run through the DataMasker-V2 pipeline before being loaded into your Jupyter environments.
- **Log Retention:** All AI inference logs must be securely retained for 5 years to comply with financial auditing standards.
- **External APIs:** Sending internal company data or customer PII to external, public LLM APIs (like public ChatGPT or standard Claude endpoints) is a critical security violation. Always use our locally hosted open-source models (e.g., Llama 3) for processing sensitive text.

4. Key Contacts & Support

If you encounter roadblocks or need elevated permissions, reach out to the following personnel:

- **Data Engineering Lead:** Sarah Wijaya (swijaya@kouventa.internal) - *Contact for database access and data pipeline failures.*
- **Security & Compliance Officer:** Budi Santoso (bsantoso@kouventa.internal) - *Contact for POJK compliance questions and API clearance.*
- **IT Helpdesk:** Ext. 4040 - *Contact for hardware issues or VPN connectivity problems.*